

Adithya Narayan

adithyanarayanb@gmail.com | adinarayanb.github.io | Boston, MA

Profile

Motivated MS CS graduate with expertise in building scalable ML platforms for image, text, and speech modeling use cases. Passionate about real-time inference pipelines, multi-modal learning, and MLOps. Possess strong collaboration/prototyping skills with over 2 years of experience across Fortune 10 and Series B setups.

Education

Northeastern University - Khoury College of Computer Sciences – Boston, MA

Master of Science in Computer Science, May 2024

Ramaiah Institute of Technology – Bangalore, India

Bachelor of Engineering in Information Science, June 2020

Experience

Machine Learning Engineer, Bear Brown LLC – Boston, MA July 2023 – Present

- Improved response relevancy by 50 %, deploying an end-to-end RAG(Retrieval Augmented Generation) LLM using Transformer embeddings obtained from 1,000 NGO-related FAQ PDFs using Python ETL pipelines.
- Reduced average user query handling time by over 70 % by deploying the fine-tuned LLM into production as a secure, scalable chatbot using React,AWS, and Flask and implementing response caching using Redis.
- Tools Used:** LangChain, LLMs, Transformer Models, Google Cloud, React, AWS, Redis, Flask

Machine Learning Research Engineer, Skit – Bangalore, Karnataka, India Dec 2021 – Aug 2022

- Increased numeric speech recognition model accuracy by over 20% for five different clients, by training DeepSpeech and hybrid CNN/GMM models, and scaling them to over 10 million customers
- Reduced deep learning model benchmarking time by 40 percent by building inference pipelines on Kubeflow
- Built internal CLI tools with Go Routines and Bash to generate high-quality synthetic speech transcriptions, enabling the solutions engineering team to fine-tune customer models in 5 hours instead of 2 days
- Increased customer resolution rate by 30% by performing A/B testing for conversation strategy selection
- Tools Used:** Pandas, Bash, Numpy, PyTorch, AWS, Kubernetes, Kubeflow, Go Routines, Metabase, PostgreSQL

SDE - Machine Learning, Stride.ai – Bangalore, Karnataka, India May 2021 – July 2021

- Developed a multilingual Named Entity Extraction microservice with an F1 score of 0.8, using Transformers, Random Forest Models, SkipGram, and RNN-LSTM, deployed across 10+ public schools
- Extended internal large dataset preprocessing platform with Airflow, authoring Python ELT pipelines to denoise and clean over 10,000 multilingual documents stored on AWS S3
- Tools Used:** Transformers, Random Forest, Pandas, SkipGram, RNN-LSTM, AWS, Flask, Kafka, Airflow

IT Software Engineer, Shell – Bangalore, Karnataka, India Sep 2020 – May 2021

- Devised and built ETL pipelines for regulatory services and reporting using SQL and Object-Oriented ABAP, querying 100+ GB daily for 500+ users across Finance and HR teams
- Tools Used:** SQL, Object-Oriented ABAP

Academic Projects

LLMs for Speech and Documents AWS, Kafka, LangChain, React, Flask

- Built and implemented a state-of-the-art LLM model using ETL pipelines incorporating RAG to transform audio and PDF files into AI-readable metadata, and created a search/natural language querying interface enabling users to interact with the LLM and locate context within the corpus.

Skills

Deep Learning: Transformers, LLM, PyTorch, TensorFlow, CNN, Seq2Seq, RNN-LSTM, YOLO-v5, Gen AI

Data Analysis: Pandas, NumPy, Scikit-Learn, Matplotlib

Backend: REST, gRPC, PostgreSQL, MongoDB, Event-Driven Microservices, Django, AWS, Redis

Orchestration/Data: Kubeflow, Kafka, Kubernetes, Airflow, Spark, Docker, Github Actions